

# Statistical Learning

## Chapter 1. Overview of Supervised learning

Pedro Delicado and Alex Sanchez

UPC & UB

# Outline

- 1 Supervised and Unsupervised learning
- 2 Statistical Decision Theory
- 3 Regression Problems
- 4 Classification problems

# Supervised Learning (the prediction problem)

- Let  $(X, Y)$  be a r.v. with support  $\mathcal{X} \times \mathcal{Y} \subseteq \mathbb{R}^p \times \mathbb{R}$ .
- General supervised learning or prediction problem:
  - Training sample:  $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$ , i.i.d. from  $(X, Y)$ .
  - The goal is to define a function (possibly depending on the sample)  $h_S : \mathcal{X} \mapsto \mathcal{Y}$  such that for a new independent observation  $(x_{n+1}, y_{n+1})$ , from which we only know  $x_{n+1}$ , it happens that:

$$\hat{y}_{n+1} = h_S(x_{n+1}) \text{ is close to } y_{n+1} \text{ (in some sense).}$$

- Function  $h_S$  is called generically a *prediction function*. (or classification function or regression function, depending on the case).

# Classification and Regression problems

The prediction function  $h_S$  is said to describe a *classification* or a *regression* problem depending on the case.

- If  $\mathcal{Y} \subseteq \mathbb{R}$  (or  $\mathcal{Y}$  an interval) we have a standard *regression problem*.
  - Example: *Relating Salary and demographic variables*
- If  $\mathcal{Y} = \{0, 1\}$  (or, also,  $\mathcal{Y} = \{-1, 1\}$ ) we have a problem of *binary classification* or discrimination.
  - Example: *Predicting if a COVID patient will require (or not) ICU*
- If  $\mathcal{Y} = \{1, \dots, K\}$  (or  $\mathcal{Y} = \{y \in \{0, 1\}^K : \sum_{k=1}^K y_k = 1\}$ ) we face a of  $K$  classes classification problem.
  - Example: *Classifying a tumor into one of many types*

# Supervised learning

- Probabilistic model for supervised learning
  - Response variable  $Y$ .
  - Explanatory variables (features)  $X = (X_1, \dots, X_p)$ .
  - Data  $(x_i = (x_{i1}, \dots, x_{ip}), y_i), i = 1, \dots, n$  i.i.d. from the random variable

$$(X = (X_1, \dots, X_p), Y) \sim \Pr(X, Y)$$

- $\Pr(X, Y)$  denotes the joint distribution of  $X$  and  $Y$ .
  - When this joint distribution is continuous,  $\Pr(X, Y)$  is the joint probability density function.

- Main interest is *predicting*  $Y$  from  $X$ .
- Given the probabilistic model it can be re-stated as *learning the conditional distribution*  $\Pr(Y \mid X)$ .
- In practice we focus on *learning a conditional location parameter*:

$$\mu(x) = \operatorname{argmin}_{\mu} \mathbb{E}(L(Y, \mu) \mid X = x),$$

where  $L(y, \hat{y})$ , loss function, measures the error of predicting  $y$  with  $\hat{y}$ .

- For quadratic loss,  $L(y, \hat{y}) = (y - \hat{y})^2$ ,  $\mu(x)$  is the regression function:

$$\mu(x) = \mathbb{E}(Y \mid X = x)$$

# Unsupervised learning

- It aims at learning relationships and structure from the observed data.
- Probabilistic model:
  - Variables of interest:  $X = (X_1, \dots, X_p)$ .
  - Data  $x_i = (x_{i1}, \dots, x_{ip})$ ,  $i = 1, \dots, n$  i.i.d. from the random variable

$$X = (X_1, \dots, X_p) \sim \Pr(X).$$

- $\Pr(X)$  denotes the probability distribution of  $X$ .
  - If  $X$  is continuous,  $\Pr(X)$  is the probability density function of  $X$ .
- Main interest: To infer properties of  $\Pr(X)$ .

# Specific problems in unsupervised learning:

- Estimating directly the density function  $\Pr(x)$  :
  - Density estimation (histogram, kernel densities, Gaussian MM)
- Detecting homogeneous subpopulations  $C_1, \dots, C_k$  s.t.:  
$$\Pr(x) = \sum_{j=1}^k \alpha_j \Pr(x | C_j), \alpha_j \geq 0, \sum_j \alpha_j = 1.$$
  - Clustering (hierarchical clustering,  $k$ -means, ...)
- Finding low-dimensional hyper-planes or hyper-surfaces (manifolds) in  $\mathbb{R}^p$  around which the probability  $\Pr(x)$  is concentrated.
  - Dimensionality reduction (PCA, MDS, Manifold learning ...)
- Proposing generative probabilistic models for  $X$ , depending on low-dimensional unobservable random variables  $F$ .
  - Extraction of latent variables (Factor Analysis, ...)



# Statistical Decision Theory

- The *prediction problem* can be written as a *decision problem* which can be casted in the setting of *Statistical Decision Theory*.
- Let  $(X, Y)$  be a r.v. with support  $\mathcal{X} \times \mathcal{Y} \subseteq \mathbb{R}^p \times \mathbb{R}$ .
- Prediction problem: To look for a prediction function  $h : \mathcal{X} \mapsto \mathcal{Y}$  such that  $h(X)$  is close to  $Y$  in some sense.
- The (lack of) closeness between  $h(X)$  and  $Y$  is usually measured by a loss function  $L(Y, h(X))$ .
  - For instance, the squared error loss is  $L(Y, h(X)) = (Y - h(X))^2$ .
  - $L(Y, h(X))$  is a r.v., with expected value  $\text{EL}(h) = \mathbb{E}(L(Y, h(X)))$ , called expected loss, that only depends on  $h$ .
- Decision problem: To find the prediction function  $h : \mathcal{X} \mapsto \mathcal{Y}$  that minimizes the expected loss.

# Bayes rule

- Denote by  $\Pr_{(X,Y)}(x,y)$  the joint probability distribution of  $(X,Y)$ .
- Observe that, for any  $h : \mathcal{X} \mapsto \mathcal{Y}$  a lower bound for  $\text{EL}(h)$  can be set as follows:

$$\begin{aligned}\text{EL}(h) &= \mathbb{E}(L(Y, h(X))) \\ &= \int_{\mathcal{X} \times \mathcal{Y}} L(y, h(x)) d\Pr_{(X,Y)}(x, y) \\ &= \int_{\mathcal{X}} \left( \int_{\mathcal{Y}} L(y, h(x)) d\Pr_{Y|X=x}(y) \right) d\Pr_X(x) \\ &= \int_{\mathcal{X}} \mathbb{E}(L(Y, h(x)) \mid X = x) d\Pr_X(x) \\ &\geq \int_{\mathcal{X}} \min_{y \in \mathcal{Y}} \mathbb{E}(L(Y, y) \mid X = x) d\Pr_X(x) \\ &= \text{EL}(h_B) .\end{aligned}$$

- From the previous bound:  $EL(h) \geq EL(h_B)$ , it follows that, given a loss function  $L(y, h(x))$  no prediction function can be better than the Bayes rule, or equivalently, that
- The optimal prediction function is the Bayes rule or Bayes classifier defined as:

$$h_B(x) = \arg \min_{y \in \mathcal{Y}} \mathbb{E}(L(Y, y) \mid X = x).$$

# The regression problem

- Let  $(X, Y)$  be a  $(p + 1)$ -dimensional random variable, with  $Y \in \mathbb{R}$ .
- The regression problem: To predict  $Y$  from known values of  $X$ .
- The most common (and convenient) approach is to adopt as loss function is the *squared error loss*:  
$$L(Y, h(X)) = (Y - h(X))^2 .$$
- Expected loss known as *Prediction Mean Squared Error*, (*PMSE*):

$$\text{PMSE}(h) = \mathbb{E} ((Y - h(X))^2) .$$

- The Bayes rule in this case is

$$h_B(x) = \arg \min_{y \in \mathcal{Y}} \mathbb{E} ((Y - y)^2 \mid X = x) .$$

- Observe that, for any  $y \in \mathcal{Y}$  one can decompose the conditional expectation of the squared deviation between  $Y$  and  $y$  given  $X = x$ ,  $\mathbb{E}((Y - y)^2 | X = x)$  in such a way that:

$$\begin{aligned}\mathbb{E}((Y - y)^2 | X = x) &= \\&= \mathbb{E}(((Y - \mathbb{E}(Y | X = x)) + (\mathbb{E}(Y | X = x) - y))^2 | X = x) \\&= \mathbb{E}((Y - \mathbb{E}(Y | X = x))^2 | X = x) + \underbrace{(\mathbb{E}(Y | X = x) - y)^2}_{\geq 0} \\&\quad + 2(\mathbb{E}(Y | X = x) - y)\mathbb{E}(Y - \mathbb{E}(Y | X = x) | X = x) \\&\geq \mathbb{E}\left((Y - \underbrace{\mathbb{E}(Y | X = x)}_{h_B(x)})^2 | X = x\right)\end{aligned}$$

# Optimal predictor in regression

- From the previous development it yields that, for regression problems, the Bayes rule is the conditional expectation of  $Y$  given  $X = x$ ,

$$h_B(x) = \mathbb{E}(Y \mid X = x),$$

- It is also known as regression function of  $Y$  over  $x$  and is usually denoted by

$$m(x) = \mathbb{E}(Y \mid X = x).$$

# Parametric regression

- Parametric regression models assume that  $m(x)$  is known except for a finite number of unknown parameters,

$$m(x) \equiv m(x; \theta), \theta \in \Theta \subseteq \mathbb{R}^q,$$

- For instance, the multiple linear regression model postulates that  $m(x) = \beta_0 + x^\top \beta_1$ , with unknown parameters  $\beta_0 \in \mathbb{R}, \beta_1 \in \mathbb{R}^p$ .
- A training sample,  $S = \{(x_1, y_1), \dots, (x_n, y_n)\}$ , i.i.d. from  $(X, Y)$ , is used to estimate the parameter  $\theta$ .
- In this case  $h_S(x) = m(x; \hat{\theta})$ , where  $\hat{\theta} = \hat{\theta}(S)$  is the estimation of  $\theta$  from sample  $S$ .

# Least squares estimation

- A usual way to estimate  $\theta$  in parametric models is by least squares:

$$\hat{\theta} = \arg \min_{\theta \in \Theta} \sum_{i=1}^n (y_i - m(x_i; \theta))^2$$

- The regression function  $m(x)$  is linear in  $x$ .
- It can be shown that, independently of the distributions,  $\hat{\theta}$  is the Best Linear Unbiased Estimator (BLUE) of  $\theta$ .
- Assuming joint normality for  $X$  and  $Y$  the LS-estimator is equivalent to the maximum likelihood estimator of  $\theta$
- In this case, the model is  $Y = m(X) + \varepsilon$ , where  $\varepsilon$  is an additive noise normally distributed with zero mean and independent from  $X$ , also normally distributed.



# Least squares estimation and prediction errors

- The LS estimator  $\hat{\theta}$  minimizes the prediction error, RSS, in the training sample.
- That is, the Residual Sum of Squares,

$$\text{RSS}(\theta) = \sum_{i=1}^n (y_i - m(x_i; \theta))^2,$$

takes its minimum value when  $\theta = \hat{\theta}$

$$\overline{\text{err}} = \text{RSS}(\hat{\theta}) = \sum_{i=1}^n (y_i - m(x_i; \hat{\theta}))^2$$

# Different types of prediction error

- $\text{RSS}(\theta)$  is the *prediction error* a theoretical quantity, based on the training sample, that needs to be estimated.
- $\overline{\text{err}}$ , known as the *training error* or the *apparent error*, is an approximation to  $\text{RSS}(\theta)$ .
- We are interested in the error associated when predicting a new observation, that is the Prediction Mean Squared Error (PMSE)

$$\text{PMSE}(\theta) = \mathbb{E} \left( (Y_{n+1} - m(x_i; \theta))^2 \right),$$

- $\overline{\text{err}}$  is an optimistic estimation of the in an observation of  $(X_{n+1}, Y_{n+1})$  independent from the training sample,

# $\overline{\text{err}}$ and PMSE are not the same

- In some cases such as in linear regression  $\overline{\text{err}}$  is a good approximation to  $\min_{\theta \in \mathbb{R}^q} \text{PMSE}(\theta)$
- But, when the parametric family  $m(x; \theta), \theta \in \Theta \subseteq \mathbb{R}^q$ , is too flexible:

$$\overline{\text{err}} < \text{PMSE}(\hat{\theta}) \neq \min_{\theta \in \mathbb{R}^q} \text{PMSE}(\theta)$$

- This is the case in non-parametric regression and in many machine learning algorithms. (Example:  $k$ -nearest neighbors regression, where the tuning parameter is  $k$  ).
- We will talk later in the course about cross-validation and tuning parameters.

# k nearest-neighbors regression

- K-NN is a flexible approach to regression or classification that, instead of relying on a *global model* based on all observations models each observation locally based on its *nearest neighbors*.
- The  $k$  nearest-neighbor estimator of  $m(t) = E(Y \mid X = t)$  is:

$$\hat{m}(t) = \frac{1}{|N_k(t)|} \sum_{i \in N_k(t)} y_i,$$

where  $N_k(t)$  is the neighborhood of  $t$  defined by the  $k$  closest points  $x_i$  in the training sample.

- Closeness is defined according to a previously chosen distance measure  $d(t, x)$ , for instance, the Euclidean distance.

# K-NN is flexible or way too flexible

- K-NN regression is a great real-world example of how model flexibility impacts training error vs. prediction error:
  - When  $k=1$  the model memorizes training data, leading to zero training error.
  - However, for a new test observation, predictions are **highly unstable** (high variance):  $\text{PMSE}(\hat{\theta}) \gg \overline{\text{err}}$ .
  - As  $k$  increases, the model becomes less flexible, reducing variance but increasing bias.
  - The optimal  $k$  balances both, minimizing PMSE
- **Conclusion:** Overly flexible models, like small  $k$  in  $k$ -NN, cause **training error to be misleading**.

# Practice session

- The R notebook `knn_regr.Rmd` illustrates the advantages and drawbacks of K-NN regression using R.
- Run along it and experiment with different settings.

# The classification problem

- Let  $(X, Y)$  be a r.v. with support  $\mathcal{X} \times \mathcal{Y} \subseteq \mathbb{R}^p \times \{1, \dots, K\}$ .
- We want to predict  $Y$  from observed values of  $X$ .
- The loss function in this case can be represented by a  $K \times K$  matrix  $L$ , that will be zero on the diagonal and nonnegative elsewhere.
  - The element  $(j, k)$  of  $L$  is  $L(j, k)$  : the price paid for classifying in class  $k$  an observation belonging to class  $j$ .

# The zero-one loss function

- A common loss function for classification is the zero-one loss function is used, where *all misclassifications are charged a single unit*.
- With the 0-1 loss function the Bayes rule is

$$\begin{aligned}h_B(x) &= \arg \min_{y \in \mathcal{Y}} \mathbb{E} (L_{0-1}(Y, y) \mid X = x) \\&= \arg \min_{k \in \{1, \dots, K\}} \sum_{j=1}^K L_{0-1}(j, k) \Pr(Y = j \mid X = x) \\&= \arg \min_{k \in \{1, \dots, K\}} (1 - \Pr(Y = k \mid X = x)) \\&= \arg \max_{k \in \{1, \dots, K\}} \Pr(Y = k \mid X = x).\end{aligned}$$

- In this context the Bayes rule is known as the Bayes classifier, and says that we classify to the most probable class, conditional to the observed value  $x$  of  $X$ .



# The problem of binary classification

- In the binary classification problem:  $\mathcal{Y} = \{0, 1\}$ . Then  $(Y \mid X = x) \sim \text{Bernoulli}(p = p(x) = \Pr(Y = 1 \mid X = x) = \mathbb{E}(Y \mid X = x))$ .

- The Bayes classifier is

$$h_B(x) = \begin{cases} 1 & \text{if } p(x) \geq 1/2 \\ 0 & \text{if } p(x) < 1/2 \end{cases}$$

- As  $p(x)$  is unknown, we use a training sample to estimate it.
- Let  $(x_1, y_1), \dots, (x_n, y_n)$  be  $n$  independent realizations of  $(X, Y)$ .
- Given an estimation  $\hat{p}(x)$  of the regression function  $p(x)$ , the estimated version of the Bayes classifier is

$$h_S(x_{n+1}) = \begin{cases} 1 & \text{if } \hat{p}(x_{n+1}) \geq 1/2 \\ 0 & \text{if } \hat{p}(x_{n+1}) < 1/2 \end{cases}$$

# Parametric estimation in binary classification

- In parametric modeling it is assumed that  $p(x) = \Pr(Y = 1 \mid X = x)$  is known except for a finite number of unknown parameters,

$$p(x) \equiv p(x; \theta), \theta \in \Theta \subseteq \mathbb{R}^q.$$

- The likelihood and log-likelihood are, respectively:

$$L(\theta) = \prod_{i=1}^n \Pr(Y_i = y_i \mid X_i = x_i) = \prod_{i=1}^n p(x_i; \theta)^{y_i} (1 - p(x_i; \theta))^{1-y_i}$$

$$\ell(\theta) = \log L(\theta) = \sum_{i=1}^n (y_i \log p(x_i; \theta) + (1 - y_i) \log (1 - p(x_i; \theta)))$$

- Let  $\hat{\theta} = \arg \max_{\theta \in \Theta} \ell(\theta)$  be the maximum likelihood estimator of  $\theta$ .
- Then  $\hat{p}(x) = p(x; \hat{\theta})$  is used to define the classification rule.

# Other optimization criteria for binary classification

- Maximum likelihood is not the only possibility for estimating  $\theta$  in  $p(x; \theta)$ .
- Alternatives:
  - Minimization of the misclassification error:

$$\hat{\theta}_{\text{Miss}} = \arg \min_{\theta \in \Theta} \sum_{i=1}^n (y_i - \mathbb{I} \{p(x_i; \theta) \geq 0.5\})^2.$$

- Least squares estimation:
$$\hat{\theta}_{LS} = \arg \min_{\theta \in \Theta} \sum_{i=1}^n (y_i - p(x_i; \theta))^2.$$
- Least absolute deviation:
$$\hat{\theta}_{LAD} = \arg \min_{\theta \in \Theta} \sum_{i=1}^n |y_i - p(x_i; \theta)|.$$
- Penalized version of these criteria, when the statistical model  $p(x; \theta)$ ,  $\theta \in \mathbb{R}^q$ , is too flexible.

# k-Nearest Neighbors (k-NN) for classification

**k-Nearest Neighbors (k-NN) is a simple and effective classification method.**

It relies on the idea that similar instances should belong to the same class.

- Given a training set ( $\mathcal{T}$ ) with labeled instances  $(x_i, y_i)$ ,
- To classify a new instance  $x$ , we:
  - 1 Find the  $k$  closest points  $x_i$  to  $x$ .
  - 2 Take the \*majority vote\$ of their corresponding labels  $y_i$ .

The decision boundary of k-NN is **nonlinear** and **flexible**, adapting to local patterns in the data.

# k-NN Classification Model

The prediction for a new point ( $x$ ) is given by:

$$\widehat{Y}(x) = \frac{1}{k} \sum_{x_i \in N_k(x)} y_i$$


where:

- $N_k(x)$  is the set of the  $k$  nearest neighbors of  $x$ .
- $y_i$  are the corresponding labels (0 or 1 in binary classification).
- Closeness is typically measured using **Euclidean distance**.

For classification:

- If  $\widehat{Y}(x) > 0.5$ , classify as **Class 1**.
- Otherwise, classify as **Class 0**.

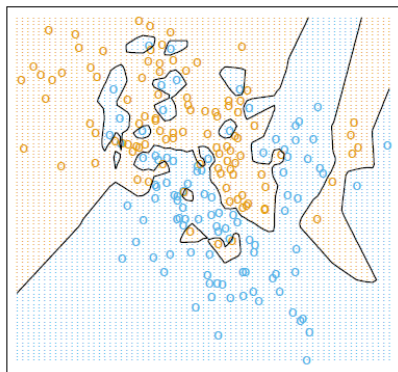
# Decision Boundaries - Linear Regression vs. k-NN



- Linear regression fits a straight line decision boundary: ( $x^T \hat{\beta} = 0.5$ ).
- Misclassifications occur because it assumes a **linear separation**.
- It does not capture **local structures** in the data.

- ◀ ◻ ▶ ◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶ ≡ 🔍 ↺

## k-NN with $k=1$



- 1-NN assigns the label of the **single closest** training point.
- Each point belongs to the nearest neighbor's class: **Voronoi tessellation**.
- Decision boundary is **highly irregular** and sensitive to noise.



# Choosing $k$ in $k$ -NN

- The parameter  $k$  in  $K$ -NN reflects its flexibility
  - $k=1$  leads to **overfitting**, that is, perfect accuracy on training but poor generalization.
  - Larger  $k$  smooths the decision boundary but might **lose fine details**.
- **Trade-offs:**
  - Small ( $k$ ) : **low bias, high variance**.
  - Large ( $k$ ) : **high bias, low variance**.
  - **Optimal  $k$**  is chosen via cross-validation (later in the course) that aims at balancing the former error measures.

# Summary

- k-NN is **flexible** and works well for complex decision boundaries.
- It is **non-parametric** and **data-driven**.
- The choice of  $(k)$  is critical for **generalization**.
- Compared to linear regression, k-NN adapts better to **nonlinear class distributions**.

# $k$ -nn classification, in R

Follow the Rmd files

`SimMixtNorm.Rmd` and `knn_class.Rmd`

# Evaluating a binary classification rule

- The explanation has been removed
- Instead you can follow the slides of a talk on Biomarkers where classification performance for binary classifiers is discussed
  - [Link to the slides](#)

# Main references